

TATA Nano RAG Diagnostics Chatbot

Technical Documentation

Presented by IntelliPredikt

Document Overview

- **Introduction:** TATA Nano RAG Diagnostics Chatbot
- **Tech Stack:** Core technologies used
- **Architecture:** High-level and detailed system design
- **System Flow:** Step-by-step process of query handling
- **Implementation Approach:** Key technical strategies
- **Data Structures & APIs:** How data is organized and exposed
- **Frontend & Security:** User interface and protection measures
- **Performance & Testing:** Optimizations and validation
- **Deployment & Decisions:** Infrastructure and rationale behind choices





Technology Stack



Backend Framework

- Flask 3.0.0: Lightweight web framework
- Flask-CORS 4.0.0: Cross-Origin support
- Gunicorn 21.2.0: WSGI HTTP server



AI & Machine Learning

- Anthropic Claude API (Sonnet 4.5): LLM
- Sentence Transformers 2.2.2: Embeddings
- scikit-learn 1.3.2: Cosine similarity
- NumPy 1.24.3: Numerical operations



Containerization

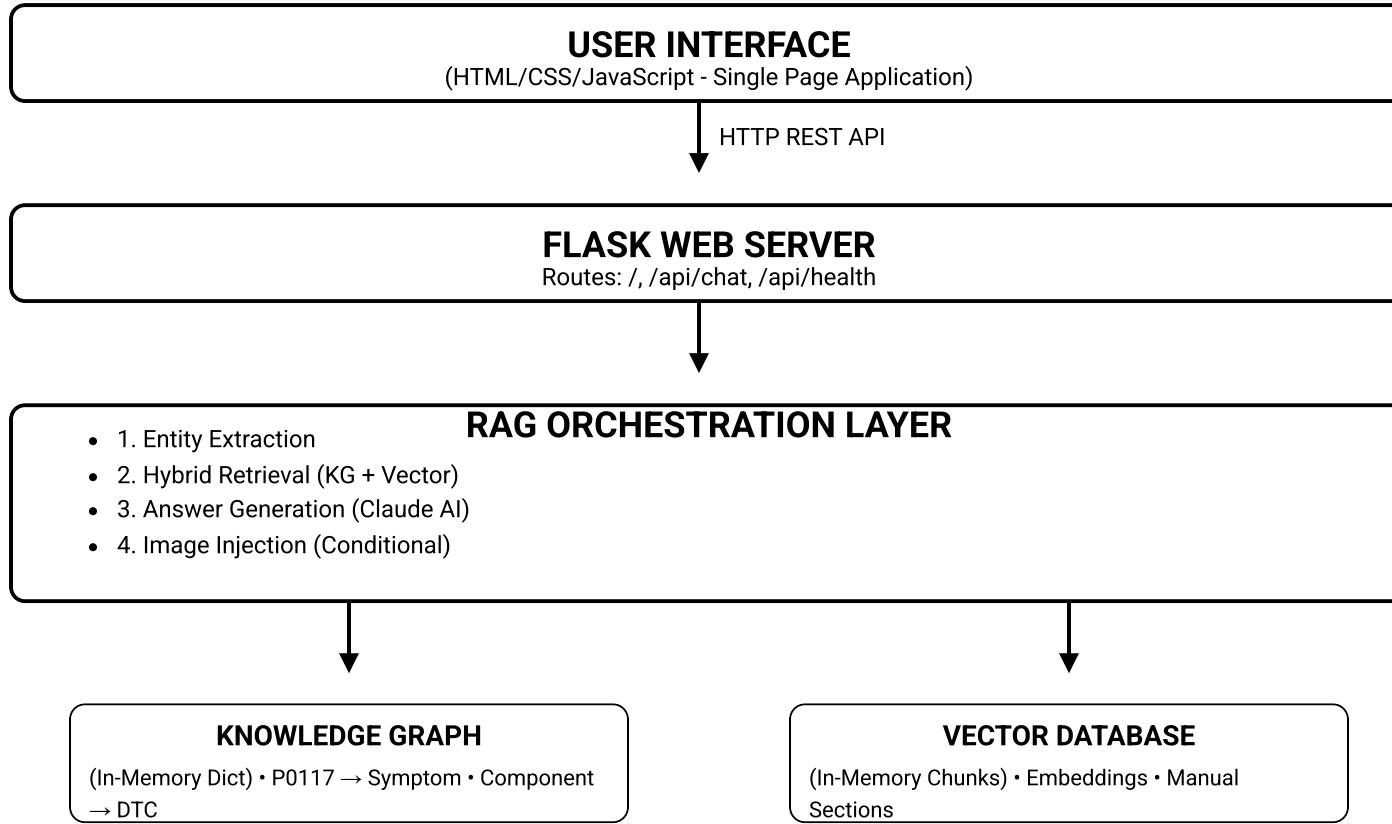
- Docker: Container platform
- Docker Compose: Orchestration
- Python 3.11-slim: Base image



Development Tools

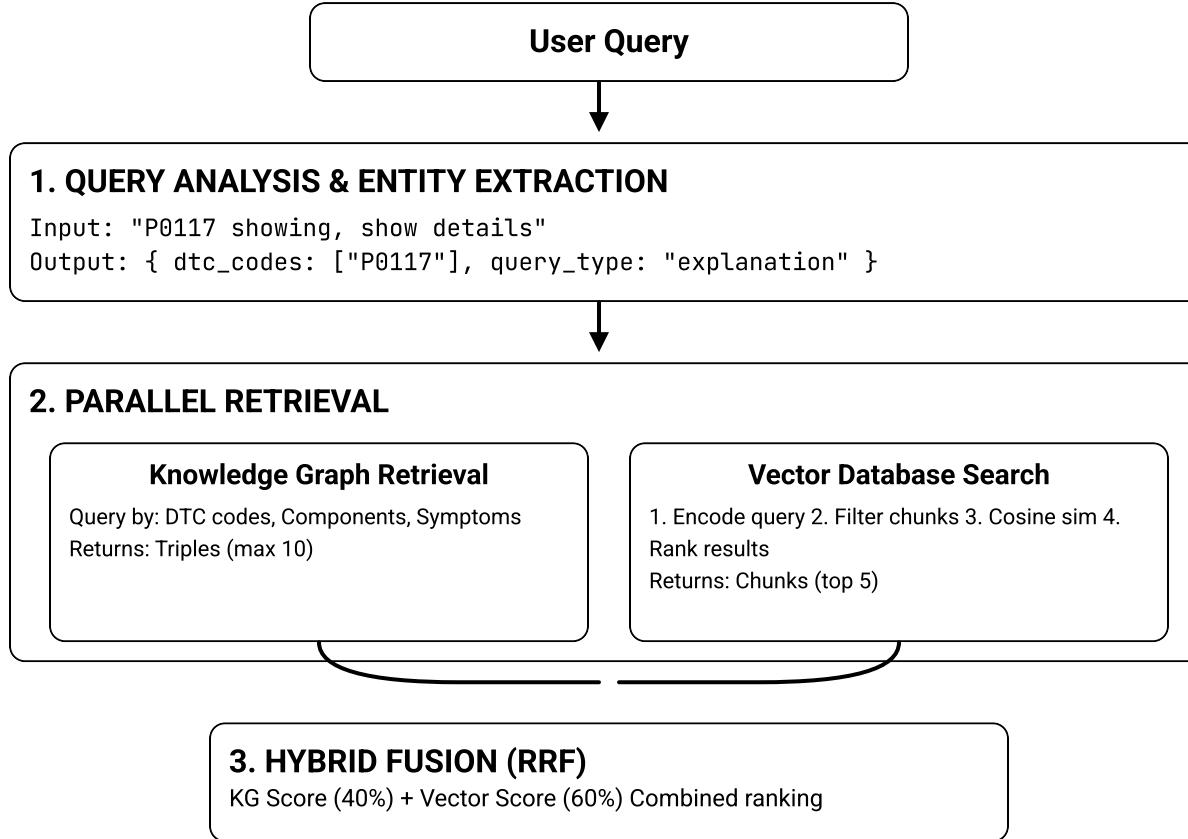
- Python 3.11: Core language
- pip: Package management

High-Level Architecture





System Flow Diagram (Part 1)





System Flow Diagram (Part 2)

4. CONTEXT PREPARATION

- Format KG triples • Format vector chunks • Build system prompt



5. CLAUDE AI GENERATION

- API Call:
- Model: claude-sonnet-4-20250514
 - System prompt + context • User query • Max tokens: 2048
- Returns: Structured HTML response



6. POST-PROCESSING

- IF query_type == "image_request": → Inject SVG diagrams
ELSE: → Return as-is



7. RESPONSE DELIVERY

JSON: { answer: "...", sources: {...} }



Implementation Approach

⌚ 1. Hybrid RAG Architecture

Why Hybrid?

- **Knowledge Graph:** Structured relationships (DTC → Symptom → Component)
- **Vector Database:** Semantic understanding of natural language
- **Combination:** Best of both worlds - precision + comprehension

Implementation: In-memory KG (dict) and Vector DB (chunks with embeddings).

🔍 2. Query Type Detection

Smart Context-Aware Responses: Query Analysis → Query Type → Response Strategy

- "P0117 details" → explanation
- "Show me picture" → image_request
- "How to fix P0117" → repair
- "Fan always running" → general

Implementation: Keyword-based entity extraction function `extract_entities(query)`.



Implementation Approach (Cont.)



3. Lazy Loading Pattern

Problem: Pre-computing embeddings causes memory issues in Docker.

Solution: Compute embeddings on first use for filtered chunks.

Benefits: Reduces startup memory, avoids crashes, caches embeddings.



4. Prompt Engineering Strategy

Context-Aware System Prompts: Dynamic prompts based on query_type.

Example: explanation prompt includes symptoms/causes, excludes repair steps.

Example: image_request prompt focuses on location/connector details.



5. Embedding Model Selection

Model: all-MiniLM-L6-v2

Why: Lightweight (80MB), Fast (~0.01s/encoding), Balanced accuracy for automotive, 384 dimensions.

Alternatives: all-mpnet-base-v2 (larger), distilbert-base-nli (lower quality).



Implementation Approach (Cont.)

6. Docker Optimization

Single Worker Architecture:

```
CMD ["gunicorn", "-w", "1", "--threads", "4", "-b", "0.0.0.0:5001", "--timeout", "120", "--preload", "chatbot_backend:app"]
```

Why: ❌ Multiple workers = SIGSEGV (model instances). ✅ Single worker + threads = shared memory. ✅ --preload forking.



7. Image Handling Strategy

Inline SVG Data URIs: COOLANT_SENSOR_IMG = ''''''

Why: ✅ No external files, ✅ Fast rendering, ✅ Scalable (vector), ✅ Small size (~2KB).

Alternatives Rejected: ❌ External image files, ❌ Base64 PNGs (larger).



Data Structures & API Endpoints



🔗 Knowledge Graph Structure

```
{ "entity_id": { "type":  
  "DTC|Component|Symptom", "relationships":  
  [...], "attributes": {...} } }  
Example: "P0117": { "type": "DTC",  
  "fault_cause": "Short Circuit to Ground",  
  "symptoms": [...] }
```

☰ Vector Chunk Structure

```
{ "id": "chunk_1", "text": "DTC P0117  
indicates...", "dtc": "P0117",  
  "component": "Coolant Sensor",  
  "embedding": np.array([...]) }
```

🔗 API Endpoints

- **GET /:** Serves frontend HTML (SPA with chat interface)
- **POST /api/chat:** Handles user queries, returns structured JSON response
- **GET /api/health:** Provides system status and configuration details



Frontend Architecture & Security Considerations



Frontend Architecture

Technology Stack:

- HTML5: Structure
- CSS3: Styling (Gradients, animations, flexbox)
- Vanilla JavaScript: No frameworks (lightweight)

Design Patterns: Async message handling with fetch API.

UI Features:

- Gradient backgrounds
- Slide-in animations
- Responsive design
- Auto-scroll messages
- Loading states



Security Considerations

1. **Environment Variables:** API keys loaded securely (fail-fast if missing).
2. **CORS Configuration:** CORS (app) to allow cross-origin requests.
3. **Input Validation:** Query length and content checked to prevent abuse.
4. **No External Dependencies:**
 - SVG images embedded
 - No CDN dependencies
 - Self-contained application



Performance, Testing & Deployment



Performance Optimizations

- Lazy Embedding Computation: Fast subsequent queries.
- In-Memory Storage: O(1) lookup, no disk I/O.
- Efficient Similarity Calculation: NumPy vectorized cosine similarity.
- Limited Context Window: Max 10 triples, 5 chunks, 2048 tokens.



Testing Approach

Query Categories Tested:

- Explanation ("P0117 details")
- Image Request ("Show me coolant sensor")
- Repair ("How to fix P0117")
- Symptom ("Fan always running")

Test Commands: curl examples for API endpoints.



Deployment Architecture & Resources

Container Structure: Gunicorn WSGI (1 worker, 4 threads, 120s timeout), Flask App, ML Models (SentenceTransformer, Embeddings cache), Data (KG & Vector chunks in-memory).

Resource Requirements:

- Memory: 2GB min (4GB recommended)
- CPU: 1 core min (2 cores recommended)
- Disk: 500MB
- Network: Internet for first startup (model download)